

Matt Mazur

[Home](#)  
[About](#)  
[Archives](#)  
[Contact](#)  
[Now](#)  
[Projects](#)

#### Follow via Email

Enter your email address to follow this blog and receive notifications of new posts by email.

Join 2,733 other followers

#### About

Hey there! I'm a data scientist at [Help Scout](#) where I wrangle data to gain insights into our product and business. I also built [Lean Domain Search](#), [Preceden](#) and [many other software products](#) over the years.



Follow me on Twitter

#### Tweets by @mhmaz

Matt Mazur Retweeted

# A Step by Step Backpropagation Example

## Background

Backpropagation is a common method for training a neural network. There is [no shortage of papers](#) online that attempt to explain how backpropagation works, but few that include an example with actual numbers. This post is my attempt to explain how it works with a concrete example that folks can compare their own calculations to in order to ensure they understand backpropagation correctly.

If this kind of thing interests you, you should [sign up for my newsletter](#) where I post about AI-related projects that I'm working on.

## Backpropagation in Python

You can play around with a Python script that I wrote that implements the backpropagation algorithm in [this Github repo](#).

## Backpropagation Visualization

For an interactive visualization showing a neural network as it learns, check out my [Neural Network visualization](#).

## Additional Resources

If you find this tutorial useful and want to continue learning about neural networks, machine learning, and deep learning, I highly recommend checking out Adrian Rosebrock's new book, [Deep Learning for Computer Vision with Python](#). I really enjoyed the book and will have a full review up soon.

## Overview

For this tutorial, we're going to use a neural network with two inputs, two hidden neurons, two output neurons. Additionally, the hidden and output neurons will include a bias.

Here's the basic structure:



**Brad Jasper**  
@bradjasper

I'm counting down the days to WWDC before figuring out the remaining plans for [@heyfocusapp](#) in 2018.

In the meantime, I've been working on a new side-project, it's all about remote work. It's called RemoteHabits.

16h

Matt Mazur Retweeted



**Naval**  
@naval

How to Get Rich (without getting lucky):

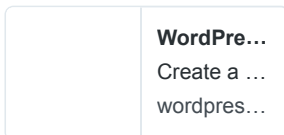
May 31, 2018



**Matt Mazur**  
@mhmazur

Replying to @mhmazur

A few years ago we A/B tested "private registration" (original) vs "privacy protection" for whois protection on [WordPress.com](#). "Privacy protection" got about 25% more people to pay for it. That small copy change has probably made hundreds of thousands of dollars by now.

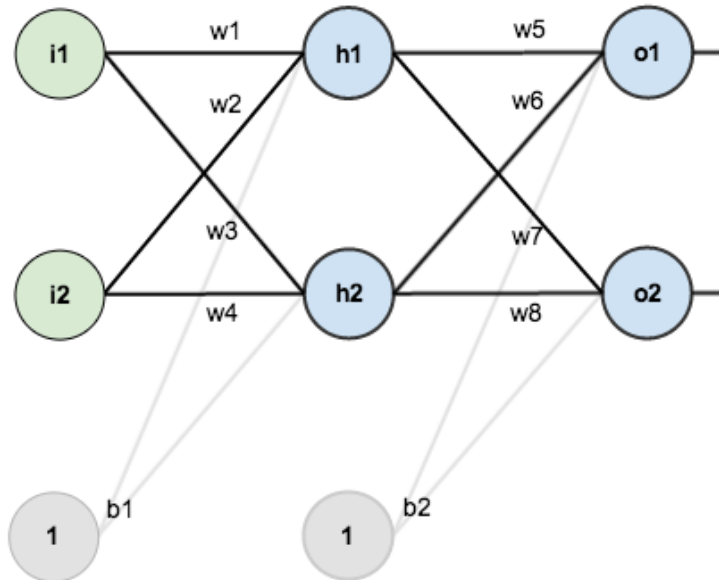


May 30, 2018

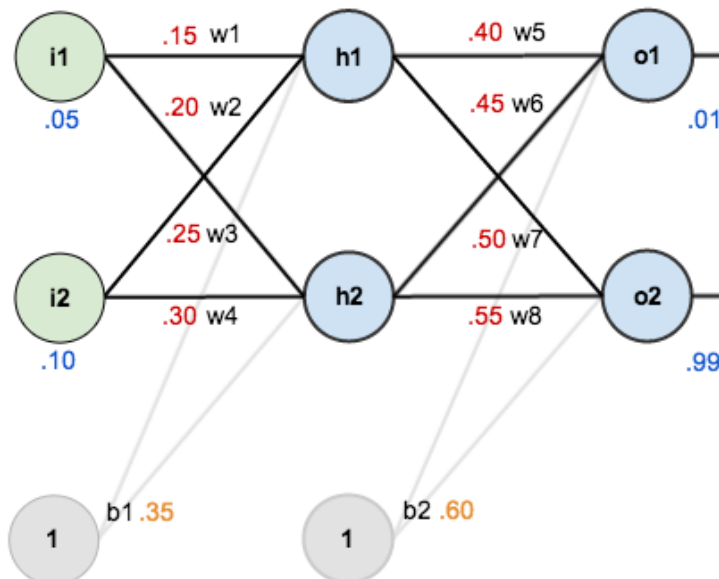


**Matt Mazur**  
@mhmazur

Love that [@Namecheap](#) is going to begin offering privacy protection for whois details for free. It's always been a clunky experience having to pay for it separately.



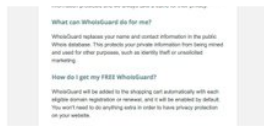
In order to have some numbers to work with, here are the **initial weights**, the **biases**, and **training inputs/outputs**:



The goal of backpropagation is to optimize the weights so that the neural network can learn how to correctly map arbitrary inputs to outputs.

For the rest of this tutorial we're going to work with a single training set: given inputs 0.05 and 0.10, we want the neural network to output 0.01 and 0.99.

## The Forward Pass



May 30, 2018



**Matt Mazur**  
@mhmazur

Replying to @mhmazur

And if salaries are actually rising within an industry, how would you distinguish that from broader participation in the salary survey? Segmenting the results by things like seniority and location can help, but might not be enough.

May 30, 2018



**Matt Mazur**  
@mhmazur

Replying to @mhmazur

It's interesting to consider though: are there industries where the average salary will go up when there's broader survey participation within the industry? Or will early adopters always tend to have higher salaries?

May 30, 2018

[Embed](#)
[View on Twitter](#)

To begin, let's see what the neural network currently predicts given the weights and biases above and inputs of 0.05 and 0.10. To do this we'll feed those inputs forward through the network.

We figure out the *total net input* to each hidden layer neuron, *squash* the total net input using an *activation function* (here we use the *logistic function*), then repeat the process with the output layer neurons.

Total net input is also referred to as just *net input* by [some sources](#).

Here's how we calculate the total net input for  $h_1$ :

$$net_{h1} = w_1 * i_1 + w_2 * i_2 + b_1 * 1$$

$$net_{h1} = 0.15 * 0.05 + 0.2 * 0.1 + 0.35 * 1 = 0.3775$$

We then squash it using the logistic function to get the output of  $h_1$ :

$$out_{h1} = \frac{1}{1+e^{-net_{h1}}} = \frac{1}{1+e^{-0.3775}} = 0.593269992$$

Carrying out the same process for  $h_2$  we get:

$$out_{h2} = 0.596884378$$

We repeat this process for the output layer neurons, using the output from the hidden layer neurons as inputs.

Here's the output for  $o_1$ :

$$net_{o1} = w_5 * out_{h1} + w_6 * out_{h2} + b_2 * 1$$

$$net_{o1} = 0.4 * 0.593269992 + 0.45 * 0.596884378 + 0.6 * 1 = 1.105905967$$

$$out_{o1} = \frac{1}{1+e^{-net_{o1}}} = \frac{1}{1+e^{-1.105905967}} = 0.75136507$$

And carrying out the same process for  $o_2$  we get:

$$out_{o2} = 0.772928465$$

### Calculating the Total Error

We can now calculate the error for each output neuron using the [squared error function](#) and sum them to get the total error:

$$E_{total} = \sum \frac{1}{2} (target - output)^2$$

[Some sources](#) refer to the target as the *ideal* and the output as the *actual*.

The  $\frac{1}{2}$  is included so that exponent is cancelled when we differentiate later on. The result is eventually multiplied by a learning rate anyway so it doesn't matter that we introduce a constant here [1].

For example, the target output for  $o_1$  is 0.01 but the neural network output 0.75136507, therefore its error is:

$$E_{o1} = \frac{1}{2}(\text{target}_{o1} - \text{out}_{o1})^2 = \frac{1}{2}(0.01 - 0.75136507)^2 = 0.274811083$$

Repeating this process for  $o_2$  (remembering that the target is 0.99) we get:

$$E_{o2} = 0.023560026$$

The total error for the neural network is the sum of these errors:

$$E_{total} = E_{o1} + E_{o2} = 0.274811083 + 0.023560026 = 0.298371109$$

## The Backwards Pass

Our goal with backpropagation is to update each of the weights in the network so that they cause the actual output to be closer the target output, thereby minimizing the error for each output neuron and the network as a whole.

### Output Layer


Consider  $w_5$ . We want to know how much a change in  $w_5$  affects the total error, aka  $\frac{\partial E_{total}}{\partial w_5}$ .

$\frac{\partial E_{total}}{\partial w_5}$  is read as “the partial derivative of  $E_{total}$  with respect to  $w_5$ ”. You can also say “the gradient with respect to  $w_5$ ”.

By applying the [chain rule](#) we know that:

$$\frac{\partial E_{total}}{\partial w_5} = \frac{\partial E_{total}}{\partial \text{out}_{o1}} * \frac{\partial \text{out}_{o1}}{\partial \text{net}_{o1}} * \frac{\partial \text{net}_{o1}}{\partial w_5}$$

Visually, here's what we're doing:



We need to figure out each piece in this equation.

First, how much does the total error change with respect to the output?

$$E_{total} = \frac{1}{2}(\text{target}_{o1} - \text{out}_{o1})^2 + \frac{1}{2}(\text{target}_{o2} - \text{out}_{o2})^2$$

$$\frac{\partial E_{total}}{\partial out_{o1}} = 2 * \frac{1}{2}(\text{target}_{o1} - \text{out}_{o1})^{2-1} * -1 + 0$$

$$\frac{\partial E_{total}}{\partial out_{o1}} = -(\text{target}_{o1} - \text{out}_{o1}) = -(0.01 - 0.75136507) = 0.74136507$$

$-(\text{target} - \text{out})$  is sometimes expressed as  $\text{out} - \text{target}$

When we take the partial derivative of the total error with respect to  $\text{out}_{o1}$ , the quantity  $\frac{1}{2}(\text{target}_{o2} - \text{out}_{o2})^2$  becomes zero because  $\text{out}_{o1}$  does not affect it which means we're taking the derivative of a constant which is zero.

Next, how much does the output of  $o_1$  change with respect to its total net input?

The partial derivative of the logistic function is the output multiplied by 1 minus the output:

$$out_{o1} = \frac{1}{1 + e^{-net_{o1}}}$$

$$\frac{\partial out_{o1}}{\partial net_{o1}} = out_{o1}(1 - out_{o1}) = 0.75136507(1 - 0.75136507) = 0.186815602$$

Finally, how much does the total net input of  $o_1$  change with respect to  $w_5$ ?

$$net_{o1} = w_5 * out_{h1} + w_6 * out_{h2} + b_2 * 1$$

$$\frac{\partial net_{o1}}{\partial w_5} = 1 * out_{h1} * w_5^{(1-1)} + 0 + 0 = out_{h1} = 0.593269992$$

Putting it all together:

$$\frac{\partial E_{total}}{\partial w_5} = \frac{\partial E_{total}}{\partial out_{o1}} * \frac{\partial out_{o1}}{\partial net_{o1}} * \frac{\partial net_{o1}}{\partial w_5}$$

$$\frac{\partial E_{total}}{\partial w_5} = 0.74136507 * 0.186815602 * 0.593269992 = 0.082167041$$

You'll often see this calculation combined in the form of the [delta rule](#):

$$\frac{\partial E_{total}}{\partial w_5} = -(target_{o1} - out_{o1}) * out_{o1}(1 - out_{o1}) * out_{h1}$$

Alternatively, we have  $\frac{\partial E_{total}}{\partial out_{o1}}$  and  $\frac{\partial out_{o1}}{\partial net_{o1}}$  which can be written as  $\frac{\partial E_{total}}{\partial net_{o1}}$ , aka  $\delta_{o1}$  (the Greek letter delta) aka the *node delta*. We can use this to rewrite the calculation above:

$$\delta_{o1} = \frac{\partial E_{total}}{\partial out_{o1}} * \frac{\partial out_{o1}}{\partial net_{o1}} = \frac{\partial E_{total}}{\partial net_{o1}}$$

$$\delta_{o1} = -(target_{o1} - out_{o1}) * out_{o1}(1 - out_{o1})$$

Therefore:

$$\frac{\partial E_{total}}{\partial w_5} = \delta_{o1} out_{h1}$$

Some sources extract the negative sign from  $\delta$  so it would be written as:

$$\frac{\partial E_{total}}{\partial w_5} = -\delta_{o1} out_{h1}$$

To decrease the error, we then subtract this value from the current weight (optionally multiplied by some learning rate, eta, which we'll set to 0.5):

$$w_5^+ = w_5 - \eta * \frac{\partial E_{total}}{\partial w_5} = 0.4 - 0.5 * 0.082167041 = 0.35891648$$

[Some sources](#) use  $\alpha$  (alpha) to represent the learning rate, [others use  \$\eta\$](#)  (eta), and [others](#) even use  $\epsilon$  (epsilon).

We can repeat this process to get the new weights  $w_6$ ,  $w_7$ , and  $w_8$ :

$$w_6^+ = 0.408666186$$

$$w_7^+ = 0.511301270$$

$$w_8^+ = 0.561370121$$

We perform the actual updates in the neural network *after* we have the new weights leading into the hidden layer neurons (ie, we use the original weights, not the updated weights, when we continue the backpropagation algorithm below).


### Hidden Layer

Next, we'll continue the backwards pass by calculating new values for  $w_1$ ,  $w_2$ ,  $w_3$ , and  $w_4$ .

Big picture, here's what we need to figure out:

$$\frac{\partial E_{total}}{\partial w_1} = \frac{\partial E_{total}}{\partial out_{h1}} * \frac{\partial out_{h1}}{\partial net_{h1}} * \frac{\partial net_{h1}}{\partial w_1}$$

Visually:



We're going to use a similar process as we did for the output layer, but slightly different to account for the fact that the output of each hidden layer neuron contributes to the output (and therefore error) of multiple output neurons. We know that  $out_{h1}$  affects both  $out_{o1}$  and  $out_{o2}$  therefore the  $\frac{\partial E_{total}}{\partial out_{h1}}$  needs to take into consideration its effect on the both output neurons:

$$\frac{\partial E_{total}}{\partial out_{h1}} = \frac{\partial E_{o1}}{\partial out_{h1}} + \frac{\partial E_{o2}}{\partial out_{h1}}$$

Starting with  $\frac{\partial E_{o1}}{\partial out_{h1}}$ :

$$\frac{\partial E_{o1}}{\partial out_{h1}} = \frac{\partial E_{o1}}{\partial net_{o1}} * \frac{\partial net_{o1}}{\partial out_{h1}}$$

We can calculate  $\frac{\partial E_{o1}}{\partial net_{o1}}$  using values we calculated earlier:

$$\frac{\partial E_{o1}}{\partial net_{o1}} = \frac{\partial E_{o1}}{\partial out_{o1}} * \frac{\partial out_{o1}}{\partial net_{o1}} = 0.74136507 * 0.186815602 = 0.138498562$$

And  $\frac{\partial net_{o1}}{\partial out_{h1}}$  is equal to  $w_5$ :

$$net_{o1} = w_5 * out_{h1} + w_6 * out_{h2} + b_2 * 1$$

$$\frac{\partial net_{o1}}{\partial out_{h1}} = w_5 = 0.40$$

Plugging them in:

$$\frac{\partial E_{o1}}{\partial out_{h1}} = \frac{\partial E_{o1}}{\partial net_{o1}} * \frac{\partial net_{o1}}{\partial out_{h1}} = 0.138498562 * 0.40 = 0.055399425$$

Following the same process for  $\frac{\partial E_{o2}}{\partial out_{h1}}$ , we get:

$$\frac{\partial E_{o2}}{\partial out_{h1}} = -0.019049119$$

Therefore:

$$\frac{\partial E_{total}}{\partial out_{h1}} = \frac{\partial E_{o1}}{\partial out_{h1}} + \frac{\partial E_{o2}}{\partial out_{h1}} = 0.055399425 + -0.019049119 = 0.036350306$$

Now that we have  $\frac{\partial E_{total}}{\partial out_{h1}}$ , we need to figure out  $\frac{\partial out_{h1}}{\partial net_{h1}}$  and then  $\frac{\partial net_{h1}}{\partial w_1}$  for each weight:

$$out_{h1} = \frac{1}{1 + e^{-net_{h1}}}$$

$$\frac{\partial out_{h1}}{\partial net_{h1}} = out_{h1}(1 - out_{h1}) = 0.59326999(1 - 0.59326999) = 0.241300709$$

We calculate the partial derivative of the total net input to  $h_1$  with respect to  $w_1$  the same as we did for the output neuron:

$$net_{h1} = w_1 * i_1 + w_3 * i_2 + b_1 * 1$$

$$\frac{\partial net_{h1}}{\partial w_1} = i_1 = 0.05$$

Putting it all together:

$$\frac{\partial E_{total}}{\partial w_1} = \frac{\partial E_{total}}{\partial out_{h1}} * \frac{\partial out_{h1}}{\partial net_{h1}} * \frac{\partial net_{h1}}{\partial w_1}$$

$$\frac{\partial E_{total}}{\partial w_1} = 0.036350306 * 0.241300709 * 0.05 = 0.000438568$$

You might also see this written as:

$$\frac{\partial E_{total}}{\partial w_1} = \left( \sum_o \frac{\partial E_{total}}{\partial out_o} * \frac{\partial out_o}{\partial net_o} * \frac{\partial net_o}{\partial out_{h1}} \right) * \frac{\partial out_{h1}}{\partial net_{h1}} * \frac{\partial net_{h1}}{\partial w_1}$$



$$\frac{\partial E_{total}}{\partial w_1} = \left( \sum_o \delta_o * w_{ho} \right) * out_{h1} (1 - out_{h1}) * i_1$$

$$\frac{\partial E_{total}}{\partial w_1} = \delta_{h1} i_1$$

We can now update  $w_1$ :

$$w_1^+ = w_1 - \eta * \frac{\partial E_{total}}{\partial w_1} = 0.15 - 0.5 * 0.000438568 = 0.149780716$$

Repeating this for  $w_2$ ,  $w_3$ , and  $w_4$

$$w_2^+ = 0.19956143$$


$$w_3^+ = 0.24975114$$

$$w_4^+ = 0.29950229$$

Finally, we've updated all of our weights! When we fed forward the 0.05 and 0.1 inputs originally, the error on the network was 0.298371109. After this first round of backpropagation, the total error is now down to 0.291027924. It might not seem like much, but after repeating this process 10,000 times, for example, the error plummets to 0.0000351085. At this point, when we feed forward 0.05 and 0.1, the two outputs neurons generate 0.015912196 (vs 0.01 target) and 0.984065734 (vs 0.99 target).

If you've made it this far and found any errors in any of the above or can think of any ways to make it clearer for future readers, don't hesitate to [drop me a note](#). Thanks!

Share this:



98 bloggers like this.

#### Related

Experimenting with a  
Neural Network-based  
Poker Bot  
In "Poker Bot"

The State of Emergent  
Mind  
In "Emergent Mind"

I'm going to write more  
often. For real this time.  
In "Writing"

Posted on [March 17, 2015](#) by [Mazur](#). This entry was posted in [Machine Learning](#) and tagged [ai](#), [backpropagation](#), [machine learning](#), [neural networks](#). Bookmark the [permalink](#).

[← Introducing  
ABTestCalculator.com, an Open  
Source A/B Test  
Significance Calculator](#)

[TetriNET Bot Source Code Published  
on Github →](#)

## 693 thoughts on “A Step by Step Backpropagation Example”

[← Older Comments](#)



**Ymi Yugy**

— February 15, 2018 at 12:56 pm

Thanks. I think I finally got a grasp on backpropagation. The only thing missing is the notation via vectors and matrices, but those shouldn't be too difficult.

[Reply](#)



**Lowell**

— February 15, 2018 at 2:41 pm

I rescind my previous agreement with Christophe. The -1 in the equation:

$$\frac{\partial E_{\text{total}}}{\partial \text{out}_{o1}} = 2 * \frac{1}{2} (\text{target}_{o1} - \text{out}_{o1})^{2-1} * -1 + 0$$

comes from the chain rule of differentiation. Which says: derivative of  $f(g)$  is

$$f'(g) * g'$$

In this case, we are taking the partial derivative with respect to  $\text{out}_{o1}$ , which makes  $f$  be  $\frac{1}{2}(g)^2$ , with  $g$  being  $(\text{target} - \text{out}_{o1})$ . So the derivative is

$$\frac{1}{2} * 2 * (\text{target} - \text{out}_{o1})^{1-1}$$

And the zero of course comes from the fact that those parts of the equation are not dependent on  $\text{out}_{o1}$  at all and are thus constant with derivative zero.

[Reply](#)

Ping!

[A 'family-tree' of words from a recurrent net – mind-builder's blog](#)



**Ben Fox**

— February 18, 2018 at 4:34 pm

Thank you! This is by far the best explanation on the topic I have seen. Is there any chance of adding a PrintFriendly/PDF link of this article? I would love to have this on my desk as a reference.

[Reply](#)

**AJ**

— February 21, 2018 at 7:28 pm

You explained so nice, thank you so much!

[Reply](#)

**Tosh Parsely**

— February 21, 2018 at 8:58 pm

How come when I run the XOR problem, the error I get never goes below  $< 0.50$ ? Even after 5000 iterations, it keeps come out to be around .50.

[Reply](#)

**Manuj Sharma**

— February 23, 2018 at 2:47 am

Very useful article. Thanks.

[Reply](#)

**Ram Sethuraman**

— February 23, 2018 at 7:22 am

Hey. I think there is a problem with this explanation.

Shouldn't  $d(E(\text{total}))/d(\text{out } h1)$  be equal to  $[d(E(\text{total}))/d(Eo1) * (d(Eo1)/d(\text{out } h1))] + [d(E(\text{total}))/d(Eo2) * (d(Eo2)/d(\text{out } h1))]$ .

But it is given to be just  $(d(Eo1)/d(\text{out } h1) + (d(Eo2)/d(\text{out } h1))$

Or am I wrong? Please help

[Reply](#)

**Ashutosh Chandra**

— April 15, 2018 at 2:29 pm

Awesome tutorial. I really understood " propagating the error backwards" concept for the first time.

[Reply](#)

**kguzel**

— February 26, 2018 at 6:37 pm

Hi,

Thank you for this great tutorial.

if it possible, can I translate your sample to Turkish on my MEDIUM essay?

Thanks a lot.

[Reply](#)



**Hans Grueber**

— February 28, 2018 at 4:41 pm

Great article, very detailed and easy to follow.

There is just something that is still unclear for me, more specifically how to calculate the new values for the biases. What bothers me is the fact that each bias is used by two neurons. For instance,  $b_1$  is used by  $h_1$  and  $h_2$ .

Then, when using the partial derivative approach to come up with “correction” to the bias, should I do it with regard to  $h_1$  or with regard to  $h_2$ ?

[Reply](#)



**David**

— March 6, 2018 at 5:24 am

Hello

I'd like to ask you when we calculate the partial derivative of  $E(\text{total})$  with respect to  $\text{out}(o_1)$ , why we have to multiply it with -1.

[Reply](#)



**student**

— March 6, 2018 at 7:17 am

I love you, writer. Thank you!

[Reply](#)



**Jonas**

— March 6, 2018 at 1:58 pm

I really don't understand why you calculated it this way:

$E/w = E/\text{out} * \text{out}/\text{net} * \text{net}/w$

instead of just calculating  $E/w$ .

$E$  and  $w$  are still used in  $E/\text{out} * \text{out}/\text{net} * \text{net}/w$  and  $\text{out}/\text{out}$  and  $\text{net}/\text{net}$  are always 1. In addition there are more limitations ( $\text{out}$  and  $\text{net}$  have to be unequal 0).

Wouldn't it be shorter and safer to just calculate  $E/w$ ?

[Reply](#)

**Michael Polkapov**

— March 10, 2018 at 11:59 pm

Thank you Very much! Really the best tutorial I have found during the search process of explanation of back propagation for dumbs. The basic has been given very Well.

[Reply](#)

Ping!

[Homer Simpson Guide to Backpropagation - Sefik Ilkin Serengil](#)**Quang Noi Truong**

— March 14, 2018 at 3:18 am

Thank for this explanation, it is helpfull for me

[Reply](#)**Satyam**

— March 15, 2018 at 12:35 pm

Can you please give an insight into momentum parameter and the math behind it with an example?

[Reply](#)**oscarjask**

— March 22, 2018 at 10:23 pm

Man this is fantastic. So good. Thank you!

[Reply](#)**Andres**

— March 27, 2018 at 10:41 am

Thank you ... your explanation was very clear ... it would be good if you translate it to code

[Reply](#)**student**

— March 29, 2018 at 6:43 am

This is a great tutorial, but I still can't figure out how to apply the backward pass on multiple hidden layers. Could you please write something on this topic? Or does anyone have a good practical example for a multiple layered NN?

[Reply](#)**Kadmian**

— April 29, 2018 at 8:25 am

All you need for backpropagation with multiple layers is here.

Try to read “next hidden layer” instead of “output layer” during the explanation of updating the hidden layer values.

Due to the chain rule, the backpropagation process is local with respect to the layers “above” and “below” the layer you are currently processing. Said another way: what is the error of the output of a hidden layer neuron? It is the weighted sum of the node deltas of the neurons in the next layer.

[Reply](#)**roie**

— May 27, 2018 at 12:41 pm

Thank you!

But what are the target values in this case. How should I replace the (output – target)?

[Reply](#)**breezegun**

— March 31, 2018 at 8:50 am

Hi, thanks for the great guide.

This really explains when there is one instance which you have one input and one target value. What if there are multiple instances with lots of input and target pairs. Then what can one do to train the network? Train for each instance and combine the weights?

Thanks

[Reply](#)**Basit**

— April 1, 2018 at 1:56 pm

Great.it cleared all my concept.thanks

[Reply](#)**Yuki**

— April 3, 2018 at 12:00 pm

Thank you for this post ! Your explanation is very clear and easy to understand !

[Reply](#)**Joy**

— April 5, 2018 at 4:18 am

It's so worthy to understand, Thanks :)

[Reply](#)

ping!

[Understanding Backpropagation. – A Machine Learning Blog](#)**goksu**

— April 8, 2018 at 4:43 am

Thank you for your information about this topic. I am confused that you updated  $w_5$ 's value, it became 0.3589... but later in the second part, you used  $w_5$ 's value as 0.40.

[Reply](#)**roie**

— May 27, 2018 at 12:36 pm

You should use the weights before the change.

[Reply](#)**Serdar**

— April 10, 2018 at 6:12 am

Thank you, now I'm able to understand what is actually going on there

[Reply](#)

ping!

[The Future of Deep Learning Research](#)**Hotaek**

— April 13, 2018 at 4:47 am

This is really fantastic

[Reply](#)**Alejandro**

— April 14, 2018 at 2:52 pm

I was struggling to find a numeric example to practise it, thank you very much

[Reply](#)



**botfactory**

— April 15, 2018 at 2:54 am

Thanks ,really help me to understand back prop

[Reply](#)

**Ashutosh Chandra**

— April 15, 2018 at 2:31 pm

Awesome tutorial. I really understood " propagating the error backwards" concept for the first time.

[Reply](#)



**Bruno**

— April 18, 2018 at 3:51 am

By testing to encode the algorithm and obtaining all  $W_n$  correct at first update, is it possible to have a error in calculation code? My second update TotalError is wrong (0.342613).

[Reply](#)



**PATRICK WONG**

— April 20, 2018 at 6:31 am

I think I've found an error for the partial  $d(out\_o1)/d(net\_o1)$ . You said  $d(out\_o1)/d(net\_o1) = out\_o1(1 - out\_o1)$  but it should be  $d(out\_o1)/d(net\_o1) = net\_o1(1 - net\_O1)$ . Other than that, great post!

[Reply](#)



**PATRICK WONG**

— April 21, 2018 at 8:17 pm

Nevermind! Silly mistake on my part (was too tired).

[Reply](#)



**Carlos**

— April 24, 2018 at 1:39 pm



## A Step by Step Backpropagation Example – Matt Mazur

I'd like to express my gratitude. I have been stuck on this Week's. And now I think I got the basis. So I can walk without crutch and continue my course. Great explanation  
Carlos valls

[Reply](#)**Pier**

— April 24, 2018 at 9:53 pm

Reblogged this on [Kernel Panic](#).[Reply](#)**eroidblog**

— April 25, 2018 at 10:35 pm

I finally understand! I remember bookmarking this post early this year or mid last year but I didn't understand it. However, after brushing up on my derivatives now I do! Thank you so much for this comprehensive tutorial!

[Reply](#)**hanifanm**

— April 25, 2018 at 11:58 pm

thanks for the post :)  
but what about the bias? dont we need to update it?

[Reply](#)**hanifanm**

— April 26, 2018 at 2:57 am

ive tried it, and my neural network can learn without updating the bias, thanks

[Reply](#)**HsintseLi**

— May 8, 2018 at 6:05 am

This link will show you the importance of the bias:

<https://stackoverflow.com/questions/2480650/role-of-bias-in-neural-networks>

[Reply](#)

**abeltre2050**

— May 10, 2018 at 2:17 pm

It can learn without the bias, but it is not practical in larger problems.  
Bias needs to be updated.

[Reply](#)**sanjay**

— April 30, 2018 at 6:23 am

its really helpful

[Reply](#)**Hamdy**

— May 3, 2018 at 6:28 am

Thank You...it so explanatory.

[Reply](#)**Robert Gonzalez (@3droberto)**

— May 4, 2018 at 5:46 pm

Hi Matt. Thanks you I was learn to implement it on GPU easily.

<https://github.com/stormcolor/gbrain>

Regards ;)

[Reply](#)**HsintseLi**

— May 8, 2018 at 6:00 am

This is exactly what 'A Step by Step Backpropagation Example' should be! Thanks for the great post!!

[Reply](#)**Monir**

— May 11, 2018 at 2:57 am

This is indeed a very good explanation . I tried to understand andrew ng 's course, which is really a brilliant course. But how does back propagation derivatives works, i had some problem in understanding . This explanation makes me feel great as now i understand it fully . Thanks Mazur.

[Reply](#)**Adewole Kayode**

— May 11, 2018 at 3:14 am

This is wonderful. Thanks for the presentation. Can we then conclude that for a backpropagation algorithm to produce optimal results, the number of iterations or epoch must be large? Is there any way around this for the network to learn faster with minimum iterations? Thanks.

[Reply](#)**Hanbit, Lee**

— May 11, 2018 at 6:53 am

I'm korean. Your explanation is best of best!  
Thank you for ginvig this story.

[Reply](#)**koushik**

— May 11, 2018 at 11:38 am

you are not updated the bias ,please ensure me that it is necessary or not.

[Reply](#)**Juel Khan**

— May 12, 2018 at 1:21 am

Your explanation is awesome. Thanks a lot.

[Reply](#)**Jerome Paddick**

— May 14, 2018 at 12:24 pm

I was really struggling with this, thanks for breaking it down and explaining it so well !

[Reply](#)**Nikhil Varma**

— May 17, 2018 at 8:05 am

Hey,

A really great explanation, thanks for sharing the knowledge.

[Reply](#)



**Dave C**

— May 24, 2018 at 10:34 am

so glad i stumbled across this site! youre my new fav blog as im a mechanical engineer shifting more towards data science/deep learning for autonomous vehicles and i enjoy playing NLHE – one stop shop for all my interests! got a lot of old posts to sift through – great work and documentation! very good simple explanations and a nice clean layout – keep up the good work!

[Reply](#)

[← Older Comments](#)

Leave a Reply

Enter your comment here...

[Blog at WordPress.com.](#)

