

Amazon Reviews Project

Use one of the following [datasets](#) to perform sentiment analysis (using Naïve Bayes and KNN, with Bag of Words, tf-idf) on the given Amazon reviews. Pick one of the "small" datasets that is a reasonable size for your computer. The goal is to create a model to algorithmically predict if a review is positive or negative just based on its text. Try to see how these reviews compare across categories. Does a review classification model for one category work for another?

Iterate by engineering new features, removing poor features, or tuning parameters. Repeat this process until you have five different versions of your classifier. Once you've iterated, answer these questions to compare the performance of each:

- Do any of your classifiers seem to overfit?
- Which seem to perform the best? Why (explain algorithmically over the dataset)?
- What features seemed to be most impactful to performance?

(*) Notes: the project should follow the guideline as previous one (Melbourne Housing)